

DCG-UPUP-Away: Automatic Symbol Learning through Grounding to Unknowns

Mycal Tucker, Derya Aksaray, Rohan Paul, and Nicholas Roy

Abstract—This paper addresses the grounding problem whose objective is to understand the meaning of natural language commands within a robot’s symbolic world. Existing techniques to solve the grounding problem typically assume that there is a fixed set of phrases or objects that the robot will encounter during deployment. However, the real world is full of confusing jargon and unique objects that are nearly impossible to anticipate and therefore train for. This paper introduces a model called the Distributed Correspondence Graph - Unknown Phrase, Unknown Percept - Away that enables to explicitly represent unknown phrases and objects as unknown, as well as reasoning about objects that are not currently perceived. The effectiveness of the model is evaluated in terms of grounding and learning new phrases via simulations and real experiments.

I. INTRODUCTION

Recently, there has been a great interest in human-robot teaming in civilian (e.g., at factories, hotels, hospitals, homes) and military (e.g., reconnaissance) applications. Communication plays an important role in effective teaming between humans and robots. One way of communication is via natural language, which provides a rich, intuitive, and flexible medium. Accordingly, the grounding problem in the literature addresses the question of how a robot can understand the meaning of a natural language command in the context of its world model (REF: G3,DCG, etc.).

The existing methods to solve the grounding problem make two primary assumptions. First, they assume a fixed set of phrases that can constitute the commands and a fixed set of objects that exist in the world model. Thus, such methods typically fail to reason about unknown phrases or objects that have never been encountered (e.g., in the training process). Second, these methods often assume that the location of the object being grounded to is known (e.g., the phrases refer to the objects that are currently perceived or localized within a known map). As a result, a robot using these methods tends to pick the most likely perceived grounding rather than exploring its surroundings.

Note that the aforementioned assumptions do not typically reflect the reality. For example, humans tend to use context-specific lexicons that the general population does not recognize. Moreover, they often refer to objects whose locations may be unknown. To deal with such cases, training a robot to know the meaning of every possible word is infeasible

All authors are with the Computer Science and Artificial Intelligence Laboratory (CSAIL), Massachusetts Institute of Technology, Cambridge, MA 02139, USA {mycal,daksaray,rohanp,nickroy}@csail.mit.edu

It appears as if acknowledgments sometimes go here, but there’s also the acknowledgment section at the end. What’s the difference?




Fig. 1: An illustration of the proposed model DCG-UPUP-Away, which has been trained only with cubes, spheres, and cylinders; and it may ground phrases to known perceived objects (1), unknown perceived objects (2), known hypothesized objects (3), or unknown hypothetical objects (4).

and inefficient. Also, attempting to reason over the space of all possible maps is similarly computationally infeasible.

This paper proposes a new model called the Distributed Correspondence Graph - Unknown Phrase, Unknown Percept - Away (DCG-UPUP-Away), which relaxes two aforementioned assumptions by 1) explicitly modeling unknown phrases and unknown percepts, and 2) creating hypothetical objects that can be out of perception. These two changes yield a model that correctly grounds a large variety of phrases in challenging environments while autonomously learning new words and symbols. This anecdotal evidence is supported by a simulation study using commands generated by Amazon Mechanical Turk users. Also, the performance of the proposed model is evaluated via real experiments where a turtlebot is initially trained to recognize a small set of phrases and objects. The results demonstrate that the robot correctly grounds commands approximately 80% of the time while learning new concepts in an unsupervised manner.

The remainder of this paper is organized as follows: The preliminaries on probabilistic graphical models used for the grounding problem is introduced in Section II. The technical approach used in developing the DCG-UPUP-Away model is presented in Section III. The model is evaluated in Section IV. Existing research in natural language robotics and human-robot interaction that complements this work are reviewed in Section V. Finally, Section VI concludes the paper by summarizing the contributions and future research.

II. BACKGROUND

The work in this paper falls within the field of natural language grounding, which addresses the problem of correctly determining how phrases relate to the real world (e.g. the phrase “go to the cube” means approaching a physical cube). In the general formulation of the grounding problem, three sets must be considered: $\Gamma = \{\text{symbolic actions and objects}\}$ is the set of groundings, which represents what phrases may be grounded to; $\Lambda = \{\text{English phrases}\}$ is the set of phrases, which represents what phrases natural language sentences may be composed of; and $\Upsilon = \Gamma^O \times \{\text{object attributes}\}$ is the cartesian product of the set of objects (i.e. $\Gamma^O \subset \Gamma$) and the object attributes (e.g. color, location), which represents the world model in which the phrases may be grounded.

Given a natural language command λ , which is a vector of phrases from the set Λ and has a length of $|\lambda|$, the general grounding problem can be formulated as a probability maximization problem

$$\gamma^* = \arg \max_{\gamma \in \Gamma^{|\lambda|}} p(\gamma | \lambda, \Upsilon), \quad (1)$$

where $\gamma \in \Gamma^{|\lambda|}$ is a vector of groundings with a length of $|\lambda|$, and Υ denotes the world model. In this formulation, the optimal vector of groundings γ^* is the one with maximum likelihood, given a command λ and a world model Υ .

In practice, the domains of Γ , Λ , and Υ in (1) typically include elements from previously seen examples. For example, rather than allowing the set of phrases Λ to include all words in a dictionary, Λ is generally assumed to only contain words that have appeared in the training examples. Moreover, solving (1) is a hard combinatorial optimization problem due to the diversity in language and world. One way to tackle this issue is to construct probabilistic graphical models based on the linguistic structure of the commands. For example, the Generalized Grounding Graph (G3) model [REF] is a factor graph that is trained from a corpus of labeled examples to ground language commands with objects, locations, and paths.

An alternative model is the Distributed Correspondence Graph (DCG) [REF], which infers the most likely set of planning constraints from language commands (rather than grounding phrases to particular actions, objects, or paths as in G3 model). In particular, DCG has the following properties: First, DCG decouples motion planning from the grounding problem. To this end, it discards the notion of grounding to specific actions or objects by instead only grounding to constraints relative to known objects (e.g. the area near a cube). Then, it passes the constraints to a motion planner. Second, DCG allows only the phrases composed entirely of the words that have already been encountered in training. Third, DCG only considers the perceived objects rather than a full model of the world including the unobserved objects. Fourth, ternary correspondence variables ϕ_{ij} are introduced to represent whether the i^{th} phrase λ_i from the overall command λ corresponds to a grounding γ_{ij} . For a given phrase λ_i , ϕ_{ij} is set to *Active* if λ_i refers to constraint γ_{ij}

(e.g., the area near the cube), *Inverted* if λ_i refers to the opposite of γ_{ij} (e.g. the area far from the cube), or *Inactive* if λ_i has no bearing on γ_{ij} (e.g. the area near a sphere). Fifth, for the sake of computational efficiency, the overall inference is factored using conditional independence according to the structure of the parse tree of λ .

Accordingly, the optimization problem solved over a DCG model becomes

$$\phi^* = \arg \max_{\phi_{ij} \in \Phi} \prod_i \prod_j p(\phi_{ij} | \gamma_{ij}, \lambda_i, \Gamma_{c_{ij}}, \Upsilon_{KP}), \quad (2)$$

where $\lambda_i \in \Lambda_{KN}$ is the i^{th} phrase in command λ and Λ_{KN} is the set of phrases with known (previously seen) words; Φ_i is the set of correspondence variables of λ_i , $\phi_{ij} \in \Phi_i$ is the j^{th} correspondence variable of λ_i ; γ_{ij} is the j^{th} grounding of λ_i ; $\Upsilon_{KP} = \{\text{object attributes}\} \times \Gamma_{KP}$ is the world model consisting of the object attributes and the set of known perceived symbolic objects Γ_{KP} ; and $\Gamma_{c_{ij}}$ is the set of child groundings of γ_{ij} . Note that $\Gamma_{c_{ij}}$ is defined as the set of groundings for the immediate children phrases (leftmost descendants) of the parent phrase λ_i in the parse tree of the natural language command.

For example, consider Fig. 2a that shows the parse tree of a simple command, and Fig. 2b that shows the corresponding DCG graphical model. The child grounding of the phrase “move” is the grounding for the phrase “to.” Similarly, the child grounding of the phrase “to” is the grounding for the phrase “the cube” (determiners such as “the” may be collapsed into their nouns). Thus, in this example, each grounding has exactly one child grounding, yielding the inter-plate structure in Fig. 2b. Examining the parse tree also reveals why the factorization in (2) is reasonable: the meaning “move” should be conditionally independent of the noun “cube” given the prepositional phrase. After all, the correct grounding of the word “move” is an action that does not depend on whether the target is a cube or a sphere, but it does depend on the position of the cube.

Finally, one must consider the factor function $\Psi : \Phi \times \Gamma \times \Lambda \times \Gamma \times \Upsilon \rightarrow \mathbb{R}$ within each plate that determines the most likely configuration of each $\phi_{ij} \in \Phi$ given $\gamma_{ij} \in \Gamma$, $\lambda_i \in \Lambda$, $\Gamma_{c_{ij}} \subset \Gamma$, and $\Upsilon_{KP} \subset \Upsilon$. Accordingly, one may rewrite (2)

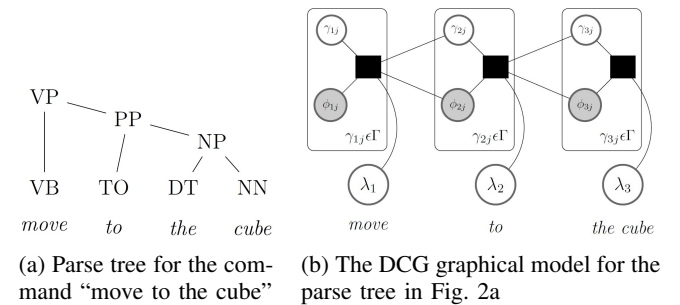


Fig. 2: An illustration of a parse tree and the corresponding DCG model.

as

$$\phi^* = \arg \max_{\phi_{ij} \in \Phi} \prod_i \prod_j^{| \Phi_i |} \Psi(\phi_{ij}, \gamma_{ij}, \lambda_i, \Gamma_{c_{ij}}, \Upsilon_{KP}), \quad (3)$$

where Ψ is a log-linear model (LLM) composed of a weighted combination of hand-coded binary functions, that is,

$$\Psi(\phi_{ij}, \gamma_{ij}, \lambda_i, \Gamma_{c_{ij}}, \Upsilon_{KP}) = \frac{\exp\left(\sum_{f \in F} \mu_f f(\phi_{ij}, \gamma_{ij}, \lambda_i, \Gamma_{c_{ij}}, \Upsilon_{KP})\right)}{\sum_{\phi_{ij} \in \{-1, 0, 1\}} \exp\left(\sum_{f \in F} \mu_f f(\phi_{ij}, \gamma_{ij}, \lambda_i, \Gamma_{c_{ij}}, \Upsilon_{KP})\right)}, \quad (4)$$

where each binary function f belongs to a set of hand-coded binary features that evaluate specific traits about a grounding (e.g., whether the word “cube” appears in λ), and μ_f is the weighting of each f . In this work, the weights μ_f are learned in a training procedure via the Limited-memory Broyden-Fletcher-Goldfarb-Shanno algorithm.

III. TECHNICAL APPROACH

In the previous section, it has been stated that the DCG model can be efficiently used in grounding problems where the objects and phrases are known and the phrases are grounded to only perceived objects. In this paper, the proposed model DCG-UPUP-Away enables the solution of a more generalized grounding problem such that 1) the phrases and objects may be known or unknown, and 2) the phrases may be grounded to objects that are out of perception. To this end, the following sections detail how to ground unknown phrases or objects, how to incrementally learn new objects and phrases, how to hypothesize groundings out of perception, and how to involve adjective attributes into grounding to avoid ambiguities.

A. Grounding Unknown Phrases or Objects

In this paper, the phrases (i.e., nouns) are defined as unknown if they have never grounded to a known object type in the training files. By maintaining a set of known nouns, one can easily determine whether or not a phrase is unknown. Similarly, the objects are defined as unknown if no object of the same class has ever appeared in the training data. Note that the DCG model does not explicitly represent the unknown symbols (i.e., objects and phrases) as in Fig. 3a, so it becomes computationally infeasible to ground them. An alternative model can be generated by decoupling the unknown symbols from the known symbols as in Fig. 3b where each known symbol is connected to each unknown symbol (as dashed edges). While such a model represents all possible correlations between the known and unknown symbols, solving the grounding problem over this graph becomes computationally expensive due to strong connectivity.

Motivated by the idea of decoupling the unknown symbols from the known ones, we propose the DCG-UPUP model which exhibits a less connected graph than the one in Fig. 3b (by removing the dashed edges). This simpler representation is mainly based on the following assumption:

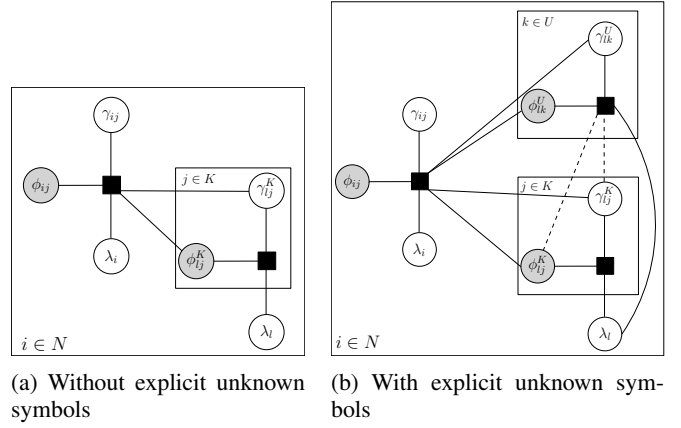


Fig. 3: Factor graph representations. Input instruction is parsed into N phrases where λ_l represents a child phrase of parent phrase λ_i . Superscripts K and U denote known and unknown variables, respectively.

given the language command λ and the world model Υ , the unknown groundings become conditionally independent from the known groundings because the variables λ and Υ are sufficient to predict the unknown groundings.

In light of going from (2) to (3), the factored objective function for the DCG-UPUP model can be written as in (5) where the domain of the world model is extended to known and unknown perceived objects (i.e., $\Upsilon_{KP} \cup \Upsilon_{UP}$).

$$\phi^* = \arg \max_{\phi_{ij} \in \Phi} \prod_i \prod_j^{| \Phi_i |} \Psi(\phi_{ij}, \gamma_{ij}, \lambda_i, \Gamma_{c_{ij}}, \Upsilon_{KP} \cup \Upsilon_{UP}), \quad (5)$$

where each feature function is an LLM model as

$$\Psi(\phi_{ij}, \gamma_{ij}, \lambda_i, \Gamma_{c_{ij}}, \Upsilon_{KP} \cup \Upsilon_{UP}) = \frac{A_U}{B_U}, \quad (6)$$

where

$$\begin{aligned} A_U &= \exp\left(\sum_{f \in F_{DCG}} \mu_f f(\phi_{ij}, \gamma_{ij}, \lambda_i, \Gamma_{c_{ij}}, \Upsilon_{KP} \cup \Upsilon_{UP})\right. \\ &\quad \left. + \sum_{f' \in F_U} \mu_{f'} f'(\phi_{ij}, \gamma_{ij}, \lambda_i, \Gamma_{c_{ij}}, \Upsilon_{KP} \cup \Upsilon_{UP})\right), \\ B_U &= \sum_{\phi_{ij} \in \{0, 1\}} \exp\left(\sum_{f \in F_{DCG}} \mu_f f(\phi_{ij}, \gamma_{ij}, \lambda_i, \Gamma_{c_{ij}}, \Upsilon_{KP} \cup \Upsilon_{UP})\right. \\ &\quad \left. + \sum_{f' \in F_U} \mu_{f'} f'(\phi_{ij}, \gamma_{ij}, \lambda_i, \Gamma_{c_{ij}}, \Upsilon_{KP} \cup \Upsilon_{UP})\right), \end{aligned}$$

F_{DCG} and F_U are the sets of hand-coded binary features used in the DCG model and for detecting unknown phrases or objects, respectively.

B. Incremental Unsupervised Learning

While the DCG-UPUP model can reason about unknown phrases and objects, it can also learn new symbols permanently. This is mainly achieved by the following unsupervised learning procedure: whenever an unknown phrase is

grounded to an unknown object, a new type is created based on the new phrase. Then, a new training example, in which the phrase grounds to the new type, is generated. Accordingly, the LLM models are retrained with the expanded set of training examples. Consequently, when the newly generated objects (or phrases) are encountered again, they become known symbols. Hence, the DCG-UPUP model can learn to associate a new phrase with a new type.

C. Hypothesized Groundings out of Perception

some intro for hypothetical grounding

$$\phi^* = \arg \max_{\phi_{ij} \in \Phi} \prod_i \prod_j \Psi(\phi_{ij}, \gamma_{ij}, \lambda_i, \Gamma_{c_{ij}}, \Upsilon_{KP} \cup \Upsilon_{UP} \cup \Upsilon_{KH} \cup \Upsilon_{UH}), \quad (7)$$

where each feature function is an LLM model as:

$$\Psi(\phi_{ij}, \gamma_{ij}, \lambda_i, \Gamma_{c_{ij}}, \Upsilon_{KP} \cup \Upsilon_{UP} \cup \Upsilon_{KH} \cup \Upsilon_{UH}) = \frac{A_{UH}}{B_{UH}} \quad (8)$$

where

$$\begin{aligned} A_{UH} &= \exp \left(\sum_{f \in F_{DCG}} \mu_f f(\phi_{ij}, \gamma_{ij}, \lambda_i, \Gamma_{c_{ij}}, \Upsilon_{KP} \cup \Upsilon_{UP} \cup \Upsilon_{KH} \cup \Upsilon_{UH}) \right. \\ &\quad + \sum_{f' \in F_U} \mu_{f'} f'(\phi_{ij}, \gamma_{ij}, \lambda_i, \Gamma_{c_{ij}}, \Upsilon_{KP} \cup \Upsilon_{UP} \cup \Upsilon_{KH} \cup \Upsilon_{UH}) \\ &\quad \left. + \sum_{f'' \in F_H} \mu_{f''} f''(\phi_{ij}, \gamma_{ij}, \lambda_i, \Gamma_{c_{ij}}, \Upsilon_{KP} \cup \Upsilon_{UP} \cup \Upsilon_{KH} \cup \Upsilon_{UH}) \right), \\ B_{UH} &= \sum_{\phi_{ij} \in \{0,1\}} \exp \left(\sum_{f \in F_{DCG}} \mu_f f(\phi_{ij}, \gamma_{ij}, \lambda_i, \Gamma_{c_{ij}}, \Upsilon_{KP} \cup \Upsilon_{UP} \cup \Upsilon_{KH} \cup \Upsilon_{UH}) \right. \\ &\quad + \sum_{f' \in F_U} \mu_{f'} f'(\phi_{ij}, \gamma_{ij}, \lambda_i, \Gamma_{c_{ij}}, \Upsilon_{KP} \cup \Upsilon_{UP} \cup \Upsilon_{KH} \cup \Upsilon_{UH}) \\ &\quad \left. + \sum_{f'' \in F_H} \mu_{f''} f''(\phi_{ij}, \gamma_{ij}, \lambda_i, \Gamma_{c_{ij}}, \Upsilon_{KP} \cup \Upsilon_{UP} \cup \Upsilon_{KH} \cup \Upsilon_{UH}) \right), \end{aligned}$$

and F_H is the set of hand-coded binary features to detect if an object is hypothetical.

-the overview of the algorithm

- final summary to compare dcg, dcg-upup, dcg-upup-away as illustrated in fig 4

D. Adjective-Attribute Heuristics

some intro for color heuristics

Note that the factored objective function in this case has exactly the same form as in (7) where each feature function $\Psi(\cdot)$ is defined as follows:

$$\Psi(\phi_{ij}, \gamma_{ij}, \lambda_i, \Gamma_{c_{ij}}, \Upsilon_{KP} \cup \Upsilon_{UP} \cup \Upsilon_{KH} \cup \Upsilon_{UH}) = \frac{A_{UHC}}{B_{UHC}} \quad (9)$$

where

$$\begin{aligned} A_{UHC} &= \exp \left(\sum_{f \in F_{DCG}} \mu_f f(\phi_{ij}, \gamma_{ij}, \lambda_i, \Gamma_{c_{ij}}, \Upsilon_{KP} \cup \Upsilon_{UP} \cup \Upsilon_{KH} \cup \Upsilon_{UH}) \right. \\ &\quad + \sum_{f' \in F_U} \mu_{f'} f'(\phi_{ij}, \gamma_{ij}, \lambda_i, \Gamma_{c_{ij}}, \Upsilon_{KP} \cup \Upsilon_{UP} \cup \Upsilon_{KH} \cup \Upsilon_{UH}) \\ &\quad + \sum_{f'' \in F_H} \mu_{f''} f''(\phi_{ij}, \gamma_{ij}, \lambda_i, \Gamma_{c_{ij}}, \Upsilon_{KP} \cup \Upsilon_{UP} \cup \Upsilon_{KH} \cup \Upsilon_{UH}) \\ &\quad \left. + \sum_{f''' \in F_C} \mu_{f'''} f'''(\phi_{ij}, \gamma_{ij}, \lambda_i, \Gamma_{c_{ij}}, \Upsilon_{KP} \cup \Upsilon_{UP} \cup \Upsilon_{KH} \cup \Upsilon_{UH}) \right), \end{aligned}$$

Algorithm 1 Grounding/Learning over DCG-UPUP-Away

```

1: procedure DCG-UPUP-AWAY
2:    $M \leftarrow$  new DCG-UPUP-Away
3:    $M.\Gamma \leftarrow$  init_groundings
4:    $M.F \leftarrow$  init_features
5:    $T \leftarrow$  init_training
6:    $M.train(M.\Gamma, M.F, T)$ 
7:   while true do
8:      $\Upsilon \leftarrow$  perceive_objects(camera,  $M.\Gamma$ )
9:      $\Upsilon \leftarrow \Upsilon +$  hypothesize_objects( $M.\Gamma$ )
10:     $\lambda \leftarrow$  get_nl_command()
11:     $[\phi^*, \gamma^*] \leftarrow M.ground(\lambda, \Upsilon)$ 
12:    if  $\gamma^*.is\_hypothesized()$  then
13:      spin_in_place()
14:    else
15:      drive_to( $\gamma^*$ )
16:     $T_u \leftarrow$  gen_unsupervised_training( $\lambda, \gamma^*, \Upsilon$ )
17:    if is_unknown( $\lambda$ ) &  $\gamma^*.is\_hypothesized()$  then
18:       $\gamma' \leftarrow$  new_grounding( $\lambda, \gamma^*, \Upsilon$ )
19:       $M.\Gamma \leftarrow M.\Gamma + \gamma'$ 
20:       $M.F \leftarrow M.F + f_{word}(\lambda[noun])$ 
21:       $M.F \leftarrow M.F + f_{obj}(\gamma^*[obj])$ 
22:       $T_u \leftarrow$  replace_unknown( $T_u, \gamma'$ )
23:     $T \leftarrow T + T_u$ 
24:     $M.train(M.\Gamma, M.F, T)$ 

```

$$\begin{aligned} B_{UHC} &= \sum_{\phi_{ij} \in \{0,1\}} \exp \left(\sum_{f \in F_{DCG}} \mu_f f(\phi_{ij}, \gamma_{ij}, \lambda_i, \Gamma_{c_{ij}}, \Upsilon_{KP} \cup \Upsilon_{UP} \cup \Upsilon_{KH} \cup \Upsilon_{UH}) \right. \\ &\quad + \sum_{f' \in F_U} \mu_{f'} f'(\phi_{ij}, \gamma_{ij}, \lambda_i, \Gamma_{c_{ij}}, \Upsilon_{KP} \cup \Upsilon_{UP} \cup \Upsilon_{KH} \cup \Upsilon_{UH}) \\ &\quad + \sum_{f'' \in F_H} \mu_{f''} f''(\phi_{ij}, \gamma_{ij}, \lambda_i, \Gamma_{c_{ij}}, \Upsilon_{KP} \cup \Upsilon_{UP} \cup \Upsilon_{KH} \cup \Upsilon_{UH}) \\ &\quad \left. + \sum_{f''' \in F_C} \mu_{f'''} f'''(\phi_{ij}, \gamma_{ij}, \lambda_i, \Gamma_{c_{ij}}, \Upsilon_{KP} \cup \Upsilon_{UP} \cup \Upsilon_{KH} \cup \Upsilon_{UH}) \right), \end{aligned}$$

and F_C is the set of hand-coded binary features for detecting color phrases or properties.

IV. EVALUATION

DCG-UPUP-Away is evaluated in two experiments. First, a simulated turtlebot within randomly generated simulated environments is given a series of user-generated natural language commands. Second, a physical turtlebot is given specific commands in a laboratory environment in order to demonstrate novel behaviors enabled by DCG-UPUP-Away. Both experiments assume a perfect object recognizer that translates raw sensor data into a world model Υ that may be used by DCG-UPUP-Away, as well as an initial set of hand-labeled training examples for training the LLM to ground cubes, spheres, and cylinders. In all trials, training DCG-UPUP-Away with 53 positive examples took less than 1 minute on a Lenovo Thinkpad X1 Carbon (TODO get more specs), and grounding a command took under 40 seconds.

A. Experimental Setup

The simulated testing environments are randomly generated in Gazebo. 10 worlds are created, and each is populated with a random collection of objects in randomized locations. There are 8 possible object types (including cubes, spheres,

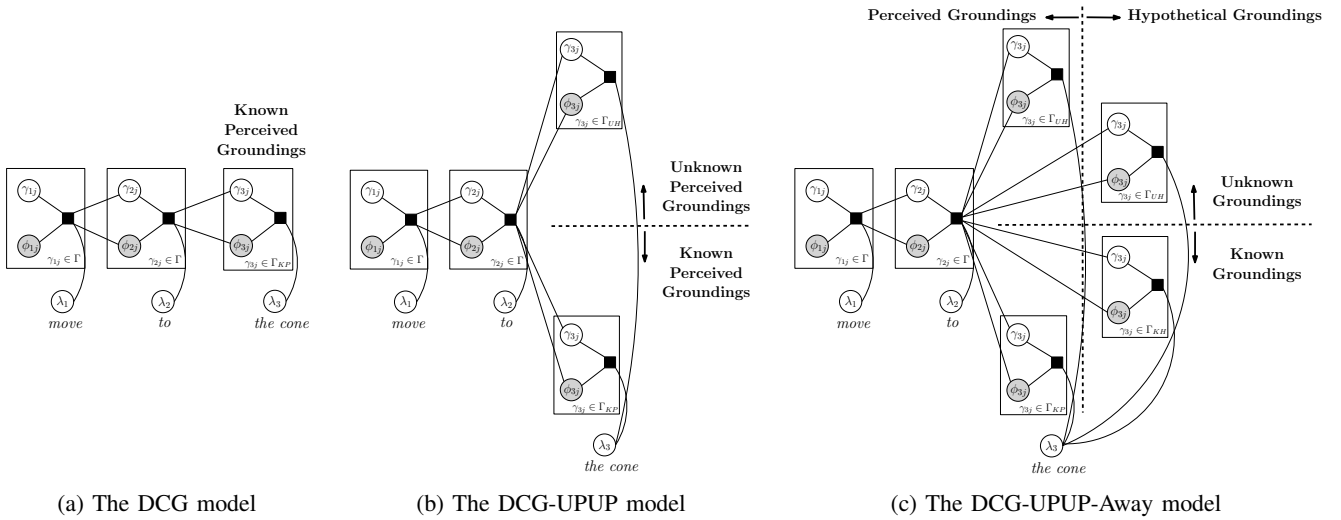


Fig. 4: The graphical models constructed for the command “move to the cone”.

and cylinders) in 3 possible colors, for a total of 24 objects. Each object has a 15% chance of being added to a given map. Using such a procedure to generate environments, when coupled with the limited field of view of the turtlebot, caused 87% of objects to be placed outside the initial field of view of the robot, demonstrating the need for the ability to ground commands to hypothesized objects.

After generating the 10 maps, screenshots of world with a single object highlighted are uploaded to Amazon Mechanical Turk. For each image, users were instructed to write a command “for approaching the highlighted object.” These image-command pairs were saved for evaluating whether a robot, when placed in the corresponding simulated world and given the natural language command, successfully approaches the correct object. An example screenshot, with an annotation supplied by a user, is shown in Figure 5.




Fig. 5: A simulated world with a highlighted object presented on Amazon Mechanical Turk, labeled by a user as “Move to the red fire hydrant.”

10 image-command pairs are randomly selected without replacement from the pool of all pairs. The 10 ordered pairs constitute one trial; each specific pair is dubbed one iteration. 30 trials are generated, each consisting of 10

iterations, for a total of 300 evaluations. When executing a trial, the turtlebot is first trained on the initial, hand-curated training set. The turtlebot is then given the natural language command from the first iteration, and then retrained using the initial data supplemented by unsupervised training examples generated by the first iteration. The retrained turtlebot is given the command from the next iteration, and appropriately retrained after each execution until all 10 iterations have been executed.

The metrics we consider are the grounding accuracy (how likely DCG-UPUP-Away is to correctly ground a phrase) and the number of known symbols. We further divide the grounding accuracy results to examine when phrases are grounded to known, unknown, or learned objects.

B. Grounding Accuracy

The primary metric used in evaluating the success of DCG-UPUP-Away is the grounding accuracy: how likely is the turtlebot to correctly execute the natural language command. Because the turtlebot is retrained between iterations, grounding accuracy may change as a function of iteration number. In fact, mean grounding accuracy remains between 70% and 90% across all iterations, as shown in Figure 6.

Although the overall grounding accuracy remains relatively constant, the underlying behavior within DCG-UPUP-Away changes over the course of a trial. In Figure 7, we plot 3 curves, showing what fraction of correctly grounded phrases refer to known objects, unknown objects, or learned objects as a function of iteration number. In the first iteration, nearly 70% of correctly grounded commands refer to known objects, but by the 10th iteration that number has fallen to nearly 10%, replaced almost entirely by correctly grounding to learned objects.

C. Learned Symbols

In order to better examine the learning behavior exhibited by DCG-UPUP-Away, the other metric considered is the

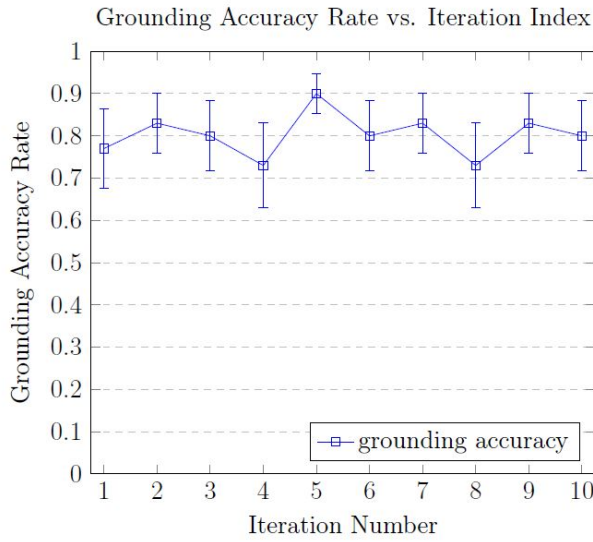


Fig. 6: this is the overall grounding accuracy




Fig. 7: this is split up (must replot well)

number of correctly known symbols. (Symbols may be incorrectly learned by associating a phrase with the wrong sort of object.) Initially, the turtlebot is trained with cubes, spheres, and cylinders, but environments may contain up to 5 additional object types (fire hydrants, drills, mailboxes, door handles, and traffic cones). Whenever an unknown phrase is grounded to such an unknown object, the turtlebot learns the new symbol. Thus, one may calculate the expected number of known symbols as a function of iteration number using combinatorics to count how many unknown objects are present. The recorded number of correctly learned symbols are plotted in Figure 8 in blue, as well as the expected number in red.

(Asking for advice: the x axis ranges from 1 to 11 because I can retrain after the 10th iteration and increase the mean number of symbols learned. It's a small increase, though, so I wouldn't be too upset cutting off the 11th iteration, though.)

As expected, the blue curve starts at 3 (for the cube, sphere, and cylinder), and stochastically monotonically increases. In 10% of trials, all 8 symbols were correctly learned; in other trials DCG-UPUP-Away incorrectly grounded unknown phrases (and therefore learned an incor-




Fig. 8: this is how we show that we learn symbols. I will update this label (and caption) of this figure. also must use std error bars instead of variance

rect symbol) or the 10 iterations collectively never referred to the five initially unknown objects, preventing DCG-UPUP-Away from ever learning the new symbol. Furthermore, learning symbols correctly improves grounding accuracy: for each additional correctly learned symbol, the turtlebot is over 4% more likely to correctly ground a command. (TODO generate p values via ANOVA.)

D. Hardware Demonstration

Outside the simulation environment, DCG-UPUP-Away was tested on a physical turtlebot in a laboratory setting. The turtlebot was placed facing a cylinder (known) and a cone (unknown). In addition, a cube (known) and a crate (unknown) were located behind the turtlebot. All objects were labeled with AR-track tags, which, in conjunction with (cite AR track package) was used in order to generate the world model Υ from a kinect mounted on the turtlebot.

Four natural language commands were used to demonstrate the full range of abilities of DCG-UPUP-Away. First, the turtlebot was given the command “move towards the cube.” The turtlebot successfully drove to the cube, demonstrating a correct grounding to a known, perceived object.

Second, the turtlebot was given the command “move towards the cone.” The turtlebot drove to the cone, demonstrating that it perceived the cone as unknown, recognized the phrase “cone” as unknown, and grounded the unknown phrase to the unknown object. Thus, a command was correctly grounded to an unknown, perceived object.

Third, the turtlebot was given the command “move towards the cube.” The turtlebot rotated in place until the cube came in view, and then approached the cube. In other words, the command was first grounded to a known, hypothesized object and then, once then cube was perceived, to a known, perceived object.

Fourth, the turtlebot was given the command “move to-

wards the crate.” Once again, the turtlebot rotated in place, this time until it saw the crate, whereupon it drove to the crate. This demonstrates two important behaviors: 1) the turtlebot must have learned what a cone was, otherwise the unknown phrase (“crate”) would have been grounded to the cone and 2) the turtlebot grounded the command to an unknown, hypothesized object until the crate was perceived. The execution of this last command, including images of the physical behavior of the turtlebot as well as the model used for grounding, is shown in Figure 9.




Fig. 9: this is a sketch of the 6 subfigures that demonstrate hypothesized groundings on hardware and in the model, and shows how it learns. I’d like this to go across the top of the page.

E. Limitations

Although the successful execution of natural language commands has been discussed in great depth, the failures of DCG-UPUP-Away are equally important in characterizing the limits of its performance.

Specifically, the most obvious limitation is that DCG-UPUP-Away assumes that, in the absence of additional information, unknown phrases refer to the first perceived unknown object. Strategies to relax this assumption by associating language adjectives with object properties have been explored in Section (TODO color section), but so far only a relatively simple method has been used.

In addition, DCG-UPUP-Away assumes a one-to-one correspondence between unknown phrases and unknown objects; it cannot, for example, learn synonyms by grounding unknown phrases to known types.

V. RELATED WORKS

related work goes here

VI. CONCLUSION

conclusion goes here. this includes future work

ACKNOWLEDGMENTS

This work was supported in part by the Robotics Consortium of the U.S Army Research Laboratory under the Collaborative Technology Alliance Program.